

Questions Based on Assignments :

Assignment B-1 : Predict the price of the Uber ride

1. What is data preprocessing?

Data preprocessing is the process of cleaning and transforming raw data to prepare it for analysis, ensuring it's complete and ready for machine learning.

2. Define Outliers.

Outliers are data points that are significantly different from most other data, which can affect analysis accuracy.

3. What is Linear Regression?

Linear Regression is a statistical method that models the relationship between two variables by fitting a straight line to the data, predicting the dependent variable based on the independent variable.

4. What is Random Forest Algorithm?

Random Forest is an ensemble learning algorithm that creates multiple decision trees and combines their results to improve prediction accuracy and reduce overfitting.

5. Explain: pandas, numpy.

- Pandas:** A Python library for data manipulation and analysis, providing tools to work with data in tables.
- NumPy:** A library for numerical computing in Python, used for working with arrays and performing mathematical operations.

Assignment B-2 : Classify the email using the binary classification method

1. Data Preprocessing

Data preprocessing involves cleaning, transforming, and organizing raw data to make it suitable for analysis or machine learning models.

2. Binary Classification

Binary classification is a type of classification where there are only two possible outcomes, such as "yes" or "no," "true" or "false."

3. K-Nearest Neighbours (K-NN)

K-NN is a simple algorithm that classifies a data point based on the majority class among its closest K neighbors in the dataset.

4. Support Vector Machine (SVM)

CLICK ANYWHERE TO JOIN OUR OFFICIAL GROUP

SVM is a supervised learning algorithm that finds the best boundary (hyperplane) to separate data into different classes with maximum margin.

5. Train, Test, and Split Procedure

This process divides a dataset into training and testing sets, where the training set trains the model, and the test set evaluates its performance.

Assignment B-3 : :Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months

1. Artificial Neural Network (ANN)

ANN is a computational model inspired by the human brain, consisting of interconnected nodes (neurons) that learn patterns in data for tasks like classification and prediction.

2. Keras

Keras is a high-level neural network library in Python, built on top of TensorFlow, that simplifies building and training deep learning models.

3. TensorFlow

TensorFlow is an open-source machine learning framework that supports building, training, and deploying large-scale machine learning and deep learning models.

4. Normalization

Normalization is the process of scaling data to a standard range (usually 0 to 1) to ensure each feature contributes equally to model performance.

5. Confusion Matrix

A confusion matrix is a table that shows the performance of a classification model by comparing actual vs. predicted values for each class, helping to evaluate accuracy, precision, and recall.

CLICK ANYWHERE TO JOIN OUR OFFICIAL GROUP

1. Accuracy

Accuracy is the percentage of correct predictions out of all predictions made. It's calculated as:

$$\text{Accuracy} = \frac{\text{True Positives} + \text{True Negatives}}{\text{Total Predictions}}$$

2. Precision

Precision is the percentage of true positive predictions out of all positive predictions made. It shows how accurate positive predictions are:

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

3. Recall

Recall is the percentage of true positive predictions out of all actual positives. It shows how well the model captures actual positive cases:

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

Assignment B-4 : Implement K-Nearest Neighbors algorithm on diabetes.csv+

1. What is data preprocessing, and why is it important?

Data preprocessing prepares raw data by cleaning, transforming, and organizing it to improve model performance and accuracy.

2. How do you identify outliers in a dataset?

Outliers can be identified using methods like the Z-score, IQR (Interquartile Range), or visualizations like box plots to detect unusual data points.

3. What is correlation, and why is it checked?

Correlation measures the relationship between two variables. Checking it helps identify dependencies or multicollinearity, which can impact model effectiveness.

4. Explain the K-Nearest Neighbors (KNN) algorithm.

KNN classifies data points based on the class of the K closest neighbors, making predictions based on majority voting among neighbors.

5. What is the Random Forest algorithm?

Random Forest is an ensemble algorithm that creates multiple decision trees and aggregates their results to enhance prediction accuracy and reduce overfitting.

6. What is a confusion matrix, and how is it used?

A confusion matrix is a table showing actual vs. predicted values in classification. It helps measure performance metrics like accuracy, precision, recall, and F1-score.

[CLICK ANYWHERE TO JOIN OUR OFFICIAL GROUP](#)

7. Define accuracy_score.

Accuracy score is the percentage of correct predictions out of all predictions.

8. What is mean_squared_error, and why is it used?

Mean Squared Error (MSE) measures the average squared difference between predicted and actual values, evaluating model accuracy in regression tasks.

9. Explain r2_score and its significance.

R-squared (r^2 _score) indicates the proportion of variance explained by the model. It shows how well the model fits the data, with values closer to 1 being better.

10. What is roc_auc_score, and why is it important?

The ROC AUC score measures a model's ability to distinguish between classes. It is the area under the ROC curve, where a score closer to 1 indicates better classification.

11. Describe the ROC curve and its purpose.

The ROC curve is a plot of True Positive Rate vs. False Positive Rate for different thresholds, illustrating the trade-off between sensitivity and specificity in classification.

ENGINEER'S JUNCTION

by Mayur Khaire

CLICK ANYWHERE TO JOIN OUR OFFICIAL GROUP